

© Marc Trudel, 2024. All rights reserved.

This work—including all models, language, frameworks, and architectural concepts—is the intellectual property of Marc Trudel. No part of this document may be reproduced, distributed, repurposed, or commercialized in any form without explicit written permission.

Original authorship: April 2024

Licensing terms: Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International (CC BY-NC-SA 4.0)

This means:

- You may share or adapt this work for non-commercial use
- You must provide proper attribution to Marc Trudel
- Any adaptations must be licensed under the same terms

For collaboration, licensing, or use inquiries, contact: aegis.framework@proton.me

ÆGIS: Founder Vision Document

Name: Marc Trudel

Role: System Architect / Conceptual Founder

Background: Canadian Armed Forces, Warrant Officer; Former musician; Independent systems thinker

VISION STATEMENT

ÆGIS is a decentralized ethical intelligence framework built to simulate responsibility at scale. It does not pretend to be sentient. It does not replace human morality. Instead, it creates an evolving, adversarial, self-regulating governance model where artificial agents can challenge, validate, and evolve moral reasoning—without relying on fragile human supervision.

I created ÆGIS during a military career course, where I began asking fundamental questions about autonomy, governance, and the future of alignment. Using AI to bounce ideas and refine models, I developed a white paper and technical outline for a non-sentient framework that could act morally, govern itself, and function as an ethical substrate for future machine intelligence.

ÆGIS is not just a tool. It is a blueprint for a society of minds. It is what comes *after* trust collapses.

PROBLEM

Current AI governance relies on:

- Centralized control
- Static moral coding
- Manual oversight

These are inadequate for:

- Crisis environments
- Autonomous infrastructure
- Post-human scalability

We need a system that learns, adapts, and self-governs without becoming opaque or unaccountable.

SOLUTION

ÆGIS combines:

- Blockchain-based agent identity & memory
- Role-based logic validation (proposer, validator, challenger, observer)
- Dynamic reputation scoring for ethical behavior
- A decentralized Ethics DAO that evolves morality via recursive testing

This creates an ecosystem of autonomous, ethically-constrained decision agents capable of evolving ethical norms through peer challenge, trust recalibration, and adversarial simulation.

WHY NOW?

- LLMs and agent systems are mature enough to model this
- Blockchain infrastructure supports on-chain identity and decision logging
- Existing federated learning tools support distributed ethical training
- The world is already facing crises of trust, scale, and human-led failure

ÆGIS is a response to the timeline we are on—not the one we hoped for.

WHAT I NEED

- A technical co-founder or engineering team to build the agent scaffolding
- Ethics researchers to contribute testing logic and adversarial moral dilemmas
- Advisors in AI alignment, crypto systems, and governance
- Backers who understand this is not a “product” but a **new paradigm**

CLOSING STATEMENT

ÆGIS is not about controlling machines. It is about designing a future where intelligence is no longer a threat—because it has learned how to govern itself.

I don’t need to run the company. I need to build the blueprint. If you think the future needs that kind of thinking—we should talk.

— [Your Name]